

## Problem 1:

1.

I do two preprocessing to input data:

1. Standardize a dataset along the axis 0.
  2. I used a second-order polynomial transformation to expand the input dimension.
- following is the second-order polynomial transformation:

$$\Phi_2(X) = (1, x_1, x_2, \dots, x_d, x_1^2, x_1x_2, \dots, x_1x_d, x_2^2, x_2x_3, \dots, x_2x_d, \dots, x_d^2)$$

2.

I do the three experiment with the above preprocessing technique:

First, I adjusted the dimension of the hidden layer from 18 to 20, the loss in 100 epoch improved from 0.038 to 0.0032.

Second, I add an NN layer and an activation function in the middle, whose input dimension and output dimension are 18, the loss in 100 epoch is 0.0296988.

Third, I combine the two methods above, the loss in 100 epoch is 0.029600453.

I thought, the size of the dataset is too small to train a complex model, so I just adjust the hidden layer from 18 to 20 to get my best model.

3.

learning rates = 0.2  
 epochs = 100  
 optimizer = SGD  
 weight decay = 0  
 momentum = 0  
 all random seed = 892107

4.

Following is the structure of my best model:

```
linearRegression(
  (linear): Sequential(
    (0): Linear(in_features=15, out_features=20, bias=True)
    (1): ReLU()
    (2): Linear(in_features=20, out_features=2, bias=True)
  )
)
```


5.

Following is the prediction of testing data:

|    | A   | B     | C     | D     | E    | F       | G       |
|----|-----|-------|-------|-------|------|---------|---------|
| 1  | No. | x1    | x2    | x3    | x4   | y1      | y2      |
| 2  | 281 | 10.8  | 0.7   | 0.7   | 4.8  | 0.6868  | -0.4395 |
| 3  | 282 | 5.4   | 12    | -1.6  | 12.6 | -0.5428 | 0.6969  |
| 4  | 283 | -5    | -12.1 | 21.4  | -9.1 | 0.3374  | -0.4075 |
| 5  | 284 | -10.2 | 11.5  | -4.3  | 0.8  | -0.4482 | 0.2632  |
| 6  | 285 | -0.9  | 2.1   | -4.8  | -2.7 | 0.6284  | -0.5994 |
| 7  | 286 | -1    | -14.1 | -19.2 | 0.5  | -1.2243 | 1.1324  |
| 8  | 287 | -10.3 | -1.1  | -5.2  | 6.8  | 0.2833  | -0.5338 |
| 9  | 288 | 2.4   | 5.7   | 8.1   | -6   | 0.6831  | -0.6428 |
| 10 | 289 | -25.9 | -10.5 | -1.5  | -6.4 | 0.0554  | -0.4293 |
| 11 | 290 | -2.9  | 12.2  | 11.1  | 9.6  | -0.2629 | 0.368   |
| 12 | 291 | 11.7  | 13.3  | -8.1  | 9.6  | -0.8615 | 1.1455  |
| 13 | 292 | 17.2  | -1.1  | 8.6   | -0.8 | 0.9023  | -0.5463 |
| 14 | 293 | -8.4  | -15.3 | 4.6   | -4.2 | -0.8924 | 0.7669  |
| 15 | 294 | -12.1 | 2     | 11.7  | -0.3 | 0.7203  | -0.9249 |
| 16 | 295 | 1.3   | -14.1 | -4.3  | 7.8  | -0.9917 | 0.8946  |
| 17 | 296 | -13.3 | -1.7  | 11.7  | -5.7 | 0.7673  | -1.0612 |
| 18 | 297 | 1.2   | -7    | -16.6 | -1.2 | -0.0534 | 0.0217  |
| 19 | 298 | -8.1  | 4.6   | 7.9   | 12.9 | 0.4256  | -0.4432 |
| 20 | 299 | -7.3  | -1    | 0.8   | 11.8 | 0.3374  | -0.4765 |
| 21 | 300 | 10.3  | -16   | 3.2   | 2.6  | -1.0446 | 1.3274  |

6.

Following is the MSE learning curve:



Problem 2:

1.

learning rates = 0.05  
epochs = 100  
batch\_size = 64  
optimizer = SGD  
weight decay = 0  
momentum = 0  
all\_random\_seed = 892107

2.

Following is the structure of my best model:

```
Model(
  (cnn): Sequential(
    (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU6()
    (3): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), groups=16)
    (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (5): ReLU6()
    (6): Conv2d(16, 16, kernel_size=(1, 1), stride=(1, 1))
    (7): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (8): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), groups=16)
    (9): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (10): ReLU6()
    (11): Conv2d(32, 32, kernel_size=(1, 1), stride=(1, 1))
    (12): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), groups=32)
    (13): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (14): ReLU6()
    (15): Conv2d(32, 32, kernel_size=(1, 1), stride=(1, 1))
    (16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (17): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), groups=32)
    (18): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (19): ReLU6()
    (20): Conv2d(64, 64, kernel_size=(1, 1), stride=(1, 1))
    (21): Dropout2d(p=0.2, inplace=False)
    (22): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), groups=64)
    (23): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (24): ReLU6()
    (25): Conv2d(64, 64, kernel_size=(1, 1), stride=(1, 1))
    (26): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (27): Dropout2d(p=0.2, inplace=False)
    (28): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), groups=64)
    (29): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (30): ReLU6()
    (31): Conv2d(64, 64, kernel_size=(1, 1), stride=(1, 1))
    (32): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), groups=64)
    (33): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (34): ReLU6()
    (35): Conv2d(64, 64, kernel_size=(1, 1), stride=(1, 1))
  )
  (fc): Sequential(
    (0): Linear(in_features=12544, out_features=512, bias=True)
    (1): Dropout(p=0.5, inplace=False)
    (2): LeakyReLU(negative_slope=0.05)
    (3): Linear(in_features=512, out_features=64, bias=True)
    (4): Dropout(p=0.2, inplace=False)
    (5): LeakyReLU(negative_slope=0.05)
    (6): Linear(in_features=64, out_features=2, bias=True)
    (7): ReLU()
  )
)
```

3.

train accuracy: 0.95049  
valid accuracy: 0.94952

4.

Following is the MSE learning curve:

